



TrashUQ

**An Edge AI and Federated Learning Platform
for Intelligent Waste Monitoring**

Technical Report and Demonstration Paper

Pol Llinàs
Bru Pallàs
Aleix Bertran
Aleix Rosinach

Universitat de Lleida
Repositories: [TrashUQ/](#) and [edge/](#)
May 14, 2026

Contents

Abstract	2
Executive Summary	2
1 Introduction	3
2 System Overview	3
3 Hardware and Edge Layer	4
4 MQTT Communication Layer	4
5 Backend Architecture	6
6 Database Layer	7
7 Frontend Dashboard	8
8 Federated Learning Layer	8
9 Edge Simulator and No-Hardware Demo	9
10 Real Model Runtime Preparation	9
11 Verification and Testing	10
11.1 Start the full stack	10
11.2 Backend health	10
11.3 Frontend proxy	10
11.4 MQTT monitor	10
11.5 Edge simulator	10
11.6 gRPC smoke test	11
11.7 Model load	11
11.8 Database verification	11
12 Results	11
13 Limitations	11
14 Future Work	12
15 Demo Script	12
16 Deployment Topology	13
17 Glossary	13
18 Logo Source	13
19 Conclusion	14

Executive statement. TrashUQ demonstrates a real edge-to-cloud telemetry and monitoring pipeline for intelligent waste classification. Edge clients publish MQTT telemetry and classification events, the backend persists the messages in PostgreSQL, the dashboard visualizes live device state and federated-learning metrics, and a gRPC coordinator exposes the initial Federated Learning control path. The no-hardware simulator is not a frontend mock: it publishes real MQTT traffic through the same backend, database, and dashboard path used by real edge devices.

Abstract

TrashUQ is an Edge AI and Federated Learning platform for intelligent waste monitoring. Edge devices, currently represented by a verified simulator and a prepared real camera/model runtime, publish status, metrics, classifications, events, logs, and help requests over MQTT topics rooted at `arduino/<device-id>/...`. The FastAPI backend subscribes to `arduino/+/#`, persists every MQTT message in PostgreSQL, and exposes a normalized dashboard state through `/api/dashboard/bootstrap`. The Next.js frontend consumes this backend bootstrap state and also subscribes to MQTT over WebSocket for live updates. In parallel, the backend exposes a gRPC Federated Learning coordinator on port 50051; `Join` and `GetGlobalModel` are implemented and verified, while `SubmitUpdate` is intentionally skipped in the edge smoke test because real local training and update generation are not implemented yet. The real model path is prepared around a TensorFlow Lite classification model, not an object detector.

Executive Summary

Capability	Status	Evidence
MQTT edge telemetry	Implemented	<code>edge/app/edge_simulator.py</code> publishes to <code>arduino/<device-id>/...</code> ; <code>TrashUQ/backend/app/mqtt_runtime.py</code> subscribes to <code>arduino/+/#</code> .
Backend persistence	Implemented	<code>mqtt_messages</code> , <code>device_status_latest</code> , <code>device_status_history</code> , and <code>coordinator_metrics</code> are created in <code>TrashUQ/backend/app/db.py</code> .
Dashboard live monitoring	Implemented	<code>/api/dashboard/bootstrap</code> plus MQTT WebSocket <code>ws://localhost:9001/mqtt</code> in <code>TrashUQ/frontend/app/page.tsx</code> .
gRPC FL Join/GetGlobalModel	Implemented	<code>TrashUQ/backend/app/fl.proto</code> , <code>grpc_server.py</code> , and <code>edge/scripts/test_fl_grpc.py</code> .
Real camera inference	Prepared	<code>edge/app/model_runner.py</code> , <code>camera_runtime.py</code> , and <code>real_edge_runtime.py</code> .
SubmitUpdate / local FL training	Pending	The edge smoke test skips <code>SubmitUpdate</code> ; no real local update payload exists yet.

1 Introduction

Waste monitoring systems need timely operational data: device availability, resource usage, inference latency, model confidence, event history, and fleet-level visibility. Sending raw video streams to a central server is expensive, fragile, and difficult to scale. TrashUQ therefore follows an Edge AI design: inference is intended to run close to the camera, and only structured telemetry is sent to the backend.

MQTT is used because it is lightweight, publish/subscribe oriented, and appropriate for edge telemetry. PostgreSQL gives the system durable, inspectable storage so the dashboard can be reconstructed after restarts. The web dashboard provides a real-time operational interface for demos and debugging. The Federated Learning layer prepares the coordination path for distributed model training without centralizing raw training data, although real local training is explicitly future work.

2 System Overview

The project is split across two repositories:

Repository	Responsibility
TrashUQ/	Docker Compose stack, Mosquitto, PostgreSQL, FastAPI backend, gRPC FL server, Next.js frontend.
edge/	MQTT edge client, no-hardware simulator, gRPC smoke tests, TensorFlow Lite model wrapper, camera/image/video runtime.

The implemented operational path is:

1. An edge client publishes MQTT telemetry.
2. Mosquitto receives MQTT traffic on 1883 and MQTT WebSocket traffic on 9001.
3. The backend subscribes to `arduino/+/#`.
4. The backend stores every MQTT message in PostgreSQL.
5. `/api/dashboard/bootstrap` returns normalized device, metric, event, classification, and FL state.
6. The dashboard renders bootstrap data and applies live MQTT WebSocket updates.
7. The edge client can contact the gRPC FL coordinator for `Join` and `GetGlobalModel`.

flowchart LR

```
user[Browser user] -->|HTTP :3000| frontend[Next.js frontend\ntrashuq-frontend]
frontend -->|/api/* proxy\nBACKEND_API_URL=http://backend:4000| backend[FastAPI backend\ntrashuq-backend]
frontend -->|MQTT over WebSocket\nws://localhost:9001/mqtt| mqtt[(Mosquitto broker\n:1883 / :9001)]

edge[Edge node\nsimulator or real runtime] -->|MQTT arduino/<device-id>/...| mqtt
mqtt -->|subscribe arduino/+/#| backend
backend -->|persist telemetry| db[(PostgreSQL\ntrashuq-db :5432)]
backend -->|dashboard bootstrap\n/api/dashboard/bootstrap| frontend

edge -->|gRPC Join / GetGlobalModel\nlocalhost:50051| fl[FL coordinator\ngRPC service :50051]
backend --> fl
fl -->|round/model state| backend

camera[Camera or image/video source] --> model[TFLite classification model\nmodels/trash_classifier.tflite]
model --> edge
```

3 Hardware and Edge Layer

The edge layer represents the device that will eventually read frames from a camera, run local inference, and publish telemetry to TrashUQ. The no-hardware simulator is already operational, while the real runtime is prepared for camera, image, or video input.

Element	Status	File
No-hardware simulator	Implemented	edge/app/edge_simulator.py
Reusable MQTT client	Implemented	edge/app/mqtt_client.py
gRPC FL smoke client	Implemented	edge/app/fl_client.py
Real model wrapper	Prepared	edge/app/model_runner.py
Camera/image/video runtime	Prepared	edge/app/camera_runtime.py, edge/app/real_edge_runtime.py
Detected model artifact	Prepared	edge/models/trash_classifier.tflite

The model discovered in the repository is TensorFlow Lite. The default labels are `cardboard`, `glass`, `paper`, and `plastic`. The normalized runtime output is classification-only:

```
[
  {
    "label": "plastic",
    "confidence": 0.91,
    "bbox": null
  }
]
```

Technical note. The detected model is a classification model, not an object detection model. TrashUQ must not be presented as producing real bounding boxes until an object detection model is introduced.

Multi-device scalability is represented by the topic structure:

```
flowchart LR
    subgraph Nodes["Arduino / Edge nodes"]
        n1["Arduino Node 01\nstatus metrics classification event logs"]
        n2["Arduino Node 02\nstatus metrics classification event logs"]
        n3["Arduino Node 03\nstatus metrics classification event logs"]
    end

    n1 -->|arduino/node-01/...| broker[(Mosquitto)]
    n2 -->|arduino/node-02/...| broker
    n3 -->|arduino/node-03/...| broker
    broker --> backend[Backend subscriber]
    backend --> db[(PostgreSQL)]
    backend --> dashboard[Dashboard]
```

4 MQTT Communication Layer

The MQTT topic root is `arduino`, configured through `MQTT_TOPIC_ROOT`. The backend subscribes to `arduino/+/#`, and the frontend subscribes to the same device topic family over MQTT WebSocket.

Topic	Purpose
<code>arduino/<device-id>/status</code>	Device operational state.

Topic	Purpose
arduino/<device-id>/metrics	Inference, resource, and FL-style metric telemetry.
arduino/<device-id>/classification	Classification result.
arduino/<device-id>/event	Structured runtime or detection events.
arduino/<device-id>/logs	Edge runtime logs.
arduino/<device-id>/help	Help or review requests.

Status payload:

```
{
  "device_id": "unoq-01",
  "online": true,
  "status": "online",
  "state": "running",
  "mode": "simulation",
  "cpu": 55,
  "ram": 39,
  "cpu_percent": 55,
  "ram_percent": 39,
  "temp": "47.0 C",
  "heartbeat": "42 ms",
  "model_version": "simulated-v1",
  "ts": "2026-05-14T12:00:00Z"
}
```

Metrics payload:

```
{
  "device_id": "unoq-01",
  "fps": 12.4,
  "inference_ms": 83.0,
  "cpu_percent": 55,
  "ram_percent": 39,
  "globalAccuracy": 82.4,
  "globalLoss": 0.31,
  "localLoss": 0.42,
  "localAccuracy": 80.1,
  "round": 3,
  "samplesTrained": 256,
  "drift": 2.1,
  "mode": "simulation",
  "ts": "2026-05-14T12:00:00Z"
}
```

Classification payload from the real runtime:

```
{
  "device_id": "unoq-01",
  "label": "plastic",
  "confidence": 0.91,
  "bbox": null,
  "source": "real_model",
  "model_version": "trash_classifier.tflite",
  "inference_ms": 83.0,
}
```

```

    "ts": "2026-05-14T12:00:00Z"
}

```

Event payload:

```

{
  "device_id": "unoq-01",
  "type": "classification_detected",
  "severity": "info",
  "message": "plastic detected with confidence 0.91",
  "label": "plastic",
  "confidence": 0.91,
  "source": "real_model",
  "ts": "2026-05-14T12:00:00Z"
}

```

Log payload:

```

{
  "device_id": "unoq-01",
  "level": "info",
  "message": "Inference started",
  "source": "real_model",
  "ts": "2026-05-14T12:00:00Z"
}

```

sequenceDiagram

```

    autonumber
    participant Edge as Edge simulator / real runtime
    participant Broker as Mosquitto MQTT broker
    participant Backend as FastAPI MQTT subscriber
    participant DB as PostgreSQL
    participant API as /api/dashboard/bootstrap
    participant UI as Next.js dashboard

    Edge->>Broker: PUBLISH arduino/unoq-01/status
    Edge->>Broker: PUBLISH arduino/unoq-01/metrics
    Edge->>Broker: PUBLISH arduino/unoq-01/classification
    Broker-->>Backend: SUB arduino/+/#
    Backend->>DB: INSERT mqtt_messages(topic, kind, device_id, payload_json, recorded_at)
    Backend->>DB: UPSERT device_status_latest / INSERT coordinator_metrics
    UI->>API: GET /api/dashboard/bootstrap
    API->>DB: Read recent status, metrics, events, classifications
    API-->>UI: devices, metricHistory, fl state, event streams
    Broker-->>UI: MQTT WebSocket live update

```

5 Backend Architecture

The backend is a FastAPI service. On startup, `TrashUQ/backend/app/main.py` initializes the database schema, starts the gRPC server, and starts the MQTT ingest runtime.

HTTP endpoints:

Endpoint	Purpose
GET /health	Health check returning <code>{"ok": true}</code> .

Endpoint	Purpose
GET /api/dashboard/bootstrap	Normalized dashboard state from PostgreSQL and FL coordinator state.
GET /api/fl/state	Minimal FL coordinator snapshot.

Real service topology:

Component	Responsibility	Port	Technology
frontend	Dashboard, API proxy, MQTT WebSocket client	3000	Next.js
backend	REST API, MQTT ingest, gRPC FL runtime	4000, 50051	FastAPI, grpcio
mqtt	MQTT and WebSocket broker	1883, 9001	Eclipse Mosquitto
db	Telemetry and history persistence	5432 host mapping by default	PostgreSQL 16
edge	Simulator and real runtime	local process	Python, OpenCV, MQTT, gRPC

The frontend API route `TrashUQ/frontend/app/api/[...path]/route.ts` proxies `/api/*` requests to `BACKEND_API_URL=http://backend:4000` inside Docker. Browser-side MQTT connects to `ws://localhost:9001/mqtt`, which is correct because the WebSocket is opened from the user's browser.

6 Database Layer

PostgreSQL provides a durable telemetry log and derived dashboard state. The schema is created in `TrashUQ/backend/app/db.py`.

Table	Purpose
<code>mqtt_messages</code>	Full record of every MQTT message received.
<code>device_status_latest</code>	Latest parsed device state.
<code>device_status_history</code>	Historical device states.
<code>coordinator_metrics</code>	Aggregated metrics derived from <code>metrics</code> payloads.

Main `mqtt_messages` fields:

Field	Meaning
<code>topic</code>	Full MQTT topic.
<code>kind</code>	<code>status</code> , <code>metrics</code> , <code>event</code> , <code>classification</code> , <code>help</code> , <code>logs</code> , or <code>other</code> .
<code>device_id</code>	Device identifier inferred from the topic.
<code>payload_text</code> / <code>payload</code>	Original payload as text.
<code>payload_json</code>	JSONB payload if parsing succeeds.
<code>recorded_at</code>	Millisecond timestamp.
<code>created_at</code>	SQL timestamp derived from <code>recorded_at</code> .

Verification query:


```
docker compose exec db psql -U trashuq -d dashboard -c "select topic, payload, created_at from mqtt_mes"
```

7 Frontend Dashboard

The dashboard consumes data through two live paths:

1. Initial and periodic bootstrap: `fetch("/api/dashboard/bootstrap")`.
2. Real-time updates: MQTT WebSocket `ws://localhost:9001/mqtt`.

The main dashboard path does not use mock devices. If real data is unavailable, the UI shows empty or waiting states instead of fabricated curves.

Dashboard area	Real source
Active Devices	<code>devices</code> from bootstrap plus live status MQTT.
Current Round	FL state and metrics payloads containing <code>round</code> .
Global Accuracy / Loss	<code>metrics.globalAccuracy</code> , <code>metrics.globalLoss</code> , and <code>metricHistory</code> .
Average CPU / RAM	Parsed status and metrics MQTT payloads.
Live Client Summary	Device state, CPU, RAM, mode, latest classification, confidence.
Event Stream	<code>event</code> , <code>logs</code> , and MQTT live updates.
Local Training Loss chart	<code>metricHistory</code> with <code>localLoss</code> / <code>globalLoss</code> .
Client Drift chart	<code>metricHistory</code> with <code>drift</code> .

flowchart LR

```
bootstrap["GET /api/dashboard/bootstrap"] --> state["React dashboard state"]
websocket["MQTT WebSocket\narduino/+status metrics event classification help logs"] --> state

state --> cards["Cards\nActive Devices, Round, Accuracy, Training State"]
state --> summary["Live Client Summary\nstatus, CPU, RAM, mode, latest classification"]
state --> charts["Charts\nLocal/global loss, accuracy, client drift"]
state --> streams["Event Stream\nEvents, logs, help, classifications"]

db["PostgreSQL mqtt_messages\nstatus/latest metrics/history"] --> bootstrap
mqtt["Mosquitto broker"] --> websocket
```

8 Federated Learning Layer

Federated Learning coordination is exposed through the real proto file at `TrashUQ/backend/app/fl.proto`.

RPC	Request fields	Response fields
Join	<code>client_id</code>	<code>ok</code> , <code>message</code> , <code>round</code> , <code>model_version</code> , <code>global_weights</code>
GetGlobalModel	<code>client_id</code>	<code>ok</code> , <code>message</code> , <code>round</code> , <code>model_version</code> , <code>global_weights</code>
SubmitUpdate	<code>client_id</code> , <code>round</code> , <code>num_samples</code> , <code>local_weights</code> , <code>local_loss</code> , <code>local_accuracy</code>	<code>ok</code> , <code>message</code> , <code>round_aggregated</code> , <code>current_round</code> , <code>model_version</code>

The current coordinator stores state in memory: online clients, current round, model version, global weights, and pending updates. The default model size is controlled by `FL_MODEL_SIZE=16`, and the default minimum number of clients per round is `FL_MIN_CLIENTS_PER_ROUND=2`.

Capability	Status
Edge Join	Works.
Edge GetGlobalModel	Works.
Edge SubmitUpdate	Intentionally skipped.
Real local training	Pending.

```
sequenceDiagram
    autonumber
    participant Edge as Edge FL client
    participant FL as Backend gRPC FL coordinator :50051
    participant State as In-memory FederatedCoordinator
    participant MQTT as MQTT event/log telemetry
```

```
Edge->>FL: Join(client_id)
FL->>State: register online client
State-->>FL: round, model_version, global_weights
FL-->>Edge: JoinResponse(ok=true)
```

```
Edge->>FL: GetGlobalModel(client_id)
FL->>State: read current global model
State-->>FL: round, model_version, global_weights
FL-->>Edge: GetGlobalModelResponse(ok=true)
Edge->>MQTT: event fl_connected
```

Note over Edge,FL: SubmitUpdate is intentionally skipped today because the edge repo does not yet prod

9 Edge Simulator and No-Hardware Demo

The simulator enables a full demonstration without Arduino hardware, a camera, or a local model runtime. It is not a UI mock. It publishes real MQTT messages that are ingested by the real broker, persisted by the real backend, and displayed by the real dashboard.

Run it with:

```
cd ~/Documents/TrashNet/edge
uv run python -m app.edge_simulator
```

The simulator publishes:

Type	Content
status	Online/offline state, CPU, RAM, temperature, heartbeat, mode.
metrics	FPS, inference latency, CPU/RAM, accuracy/loss, round, samples, drift.
classification	Simulated labels such as trash , plastic , paper , metal , organic , clean .
event	Detection and runtime events.
logs	Runtime loop activity.
help	Review/help requests when confidence is low.

10 Real Model Runtime Preparation

The `edge/` repository is prepared for real inference:

File	Role
edge/app/model_runner.py	Loads models/trash_classifier.tflite through the existing classifier path.
edge/app/camera_runtime.py	Provides camera, image, and video frame sources.
edge/app/real_edge_runtime.py	Runs inference, publishes MQTT payloads, and performs non-blocking FL smoke checks.
edge/scripts/test_model_load.py	Verifies model loading.
edge/scripts/test_camera_open.py	Verifies camera access.
edge/scripts/test_single_image_inference.py	Runs inference on one image.

Runtime variables:

```
EDGE_INPUT_SOURCE=camera
EDGE_CAMERA_INDEX=0
EDGE_IMAGE_PATH=
EDGE_VIDEO_PATH=
EDGE_MODEL_PATH=models/trash_classifier.tflite
EDGE_CONFIDENCE_THRESHOLD=0.5
EDGE_INFERENCE_INTERVAL_SEC=1
EDGE_MAX_FPS=10
```

Current limitation. Real model execution requires `tflite_runtime` or TensorFlow Lite support in the target Python environment. If the dependency is missing, `test_model_load.py` fails clearly instead of pretending inference works.

11 Verification and Testing

11.1 Start the full stack

```
cd ~/Documents/TrashNet/TrashUQ
docker compose down
docker compose up --build
```

11.2 Backend health

```
curl http://localhost:4000/health
curl http://localhost:4000/api/dashboard/bootstrap
```

11.3 Frontend proxy

```
curl http://localhost:3000/api/dashboard/bootstrap
docker compose exec frontend sh -lc 'wget -q0- http://backend:4000/health || curl -s http://backend:4000/health'
```

11.4 MQTT monitor

```
cd ~/Documents/TrashNet/TrashUQ
docker compose exec mqtt mosquitto_sub -h localhost -p 1883 -t 'arduino/+/' -v
```

11.5 Edge simulator

```
cd ~/Documents/TrashNet/edge
uv run python -m app.edge_simulator
```

11.6 gRPC smoke test

```
cd ~/Documents/TrashNet/edge
uv run python scripts/test_fl_grpc.py
```

Expected result:

```
Join: ok=True ...
GetGlobalModel: ok=True ...
SubmitUpdate skipped: real local training/model update is not implemented in the edge repo yet.
```

11.7 Model load

```
cd ~/Documents/TrashNet/edge
uv run python scripts/test_model_load.py
```

11.8 Database verification

```
cd ~/Documents/TrashNet/TrashUQ
docker compose exec db psql -U trashuq -d dashboard -c "select topic, payload, created_at from mqtt_mes"
```

12 Results

Result	Status
MQTT end-to-end from edge to broker	Verified.
Backend persistence in PostgreSQL	Verified.
Dashboard receives real backend/MQTT data	Verified.
Edge simulator populates dashboard cards, streams, and charts through MQTT	Implemented.
Charts use real <code>metricHistory</code> or waiting states	Implemented.
gRPC <code>Join</code> and <code>GetGlobalModel</code>	Verified.
TensorFlow Lite model wrapper	Implemented/prepared.
Camera/image/video runtime	Implemented/prepared.

Not claimed as completed:

Item	Reason
Real federated local training	No implementation currently produces <code>local_weights</code> , <code>local_loss</code> , and <code>local_accuracy</code> .
Real <code>SubmitUpdate</code>	Intentionally skipped until local training exists.
Real bounding boxes	The detected model is classification-only.
Guaranteed TFLite execution on every host	Depends on a compatible TFLite runtime installation.

13 Limitations

- Real `SubmitUpdate` is not implemented in the edge repository.
- No privacy budget tracking is implemented yet.
- There is no dataset-quality dashboard yet.
- Communication cost per FL round is not measured yet.

- The current model is classification-only, not object detection.
- Camera execution depends on target hardware, camera permissions, and TFLite runtime availability.
- The FL coordinator stores its state in memory; rounds and global weights are not persisted yet.
- MQTT/gRPC security hardening remains future work.

14 Future Work

- Connect Arduino/camera hardware.
- Install `tflite_runtime` or TensorFlow Lite on the target edge device.
- Run a real classification stream through the dashboard.
- Add real local training or fine-tuning.
- Implement `SubmitUpdate` with valid model updates.
- Persist FL rounds, global weights, and per-client metrics.
- Add a model registry and operational model versioning.
- Add per-device performance analytics.
- Add dataset and model drift detection.
- Harden security with MQTT authentication, TLS, device authorization, and gRPC credentials.
- Deploy to a real edge fleet.

15 Demo Script

Recommended live demo flow:

1. Start the platform:

```
cd ~/Documents/TrashNet/TrashUQ
docker compose up --build
```

2. Open the dashboard:

```
http://localhost:3000
```

3. Run the edge simulator:

```
cd ~/Documents/TrashNet/edge
uv run python -m app.edge_simulator
```

4. Show in the dashboard:

- `unoq-01` online.
- CPU/RAM/heartbeat updating.
- Live classification events.
- Real MQTT event stream.
- Loss/drift charts if metrics with `localLoss`, `globalLoss`, and `drift` are being published.

5. Verify database persistence:

```
cd ~/Documents/TrashNet/TrashUQ
docker compose exec db psql -U trashuq -d dashboard -c "select topic, payload, created_at from mqtt_mes"
```

6. Verify gRPC:

```
cd ~/Documents/TrashNet/edge
uv run python scripts/test_fl_grpc.py
```

7. Explain real camera/model readiness:

```
cd ~/Documents/TrashNet/edge
EDGE_CAMERA_INDEX=0 uv run python scripts/test_camera_open.py
uv run python scripts/test_model_load.py
```

16 Deployment Topology

flowchart TB

```
subgraph Host["Developer / demo host"]
  subgraph Compose["Docker Compose: TrashUQ"]
    frontend["frontend\nNext.js\n0.0.0.0:3000 -> 3000"]
    backend["backend\nFastAPI + gRPC\n0.0.0.0:4000 -> 4000\n0.0.0.0:50051 -> 50051"]
    mqtt["mqtt\nEclipse Mosquitto\n0.0.0.0:1883 -> 1883\n0.0.0.0:9001 -> 9001"]
    db["db\nPostgreSQL 16\n0.0.0.0:5432 -> 5432"]
  end
  edgeLocal["edge repo\nuv run python -m app.edge_simulator\nor app.real_edge_runtime"]
end

external["External edge device\nMQTT_HOST=<SERVER_LAN_IP>\nFL_GRPC_HOST=<SERVER_LAN_IP>"]
browser["Browser\nhttp://localhost:3000"]

browser --> frontend
frontend --> |internal HTTP| backend
frontend --> |browser WebSocket| mqtt
backend --> db
backend --> mqtt
edgeLocal --> mqtt
edgeLocal --> backend
external --> mqtt
external --> backend
```

17 Glossary

Term	Definition
Edge AI	AI inference executed close to the sensor or device.
MQTT	Lightweight publish/subscribe protocol for telemetry.
gRPC	Binary RPC framework based on HTTP/2 and Protocol Buffers.
Federated Learning	Distributed learning where clients send updates instead of raw data.
Global model	Aggregated model maintained by the FL coordinator.
Local update	Locally trained client update submitted to the coordinator.
Dashboard bootstrap	Initial backend state used to reconstruct the dashboard.
Telemetry	Operational status, metrics, events, logs, and classifications.
TFLite	TensorFlow Lite, a model format/runtime optimized for edge environments.

18 Logo Source

The cover includes the Universitat de Lleida logo from Wikimedia Commons: [File:Logo Universitat de Lleida.svg](#). The file page states that the source is www.udl.cat and identifies the asset as a

public-domain text/logo. A local copy is stored in `docs/paper/assets/udl_logo.svg` and rendered to `docs/paper/assets/udl_logo.png` for PDF generation.

19 Conclusion

TrashUQ already demonstrates a real edge-to-cloud monitoring platform: edge clients publish MQTT telemetry, Mosquitto routes it, FastAPI persists it in PostgreSQL, and the Next.js dashboard visualizes live device and metric state. The simulator enables a complete no-hardware demo without faking frontend data. The gRPC path validates the first Federated Learning coordinator operations through `Join` and `GetGlobalModel`, and the TensorFlow Lite runtime prepares the system for real camera-based classification. The next technical milestone is to close the local training loop: generate valid edge updates, call `SubmitUpdate`, persist FL rounds, and evaluate model performance across real devices.